

EDA, FE and Classification Model (Census Income Dataset)

Aryan Shriva

Linkedin: <https://www.linkedin.com/in/aryan-shriva-0811/> (<https://www.linkedin.com/in/aryan-shriva-0811/>).

GitHub <https://github.com/AryanShriva/Machine-Learning> (<https://github.com/AryanShriva/Machine-Learning>).

For the code please check out my Machine Learning repository on GitHub

1. EDA and FE

1. Data Profiling
2. Stastical analysis
3. Graphical Analysis
4. Data Cleaning
5. Data Scaling

2. Logistic Regression Model

1. Linear Regression Model
2. Performance metrics for above model
3. Hyper-Parameter Tuning for above model

3. Support Vector Classifier Model

1. Support Vector Classifier Model
2. Performance metrics for above model
3. Hyper-Parameter Tuning for above model

Dataset: <https://archive.ics.uci.edu/ml/datasets/Census+Income> (<https://archive.ics.uci.edu/ml/datasets/Census+Income>).

1.0 Importing required libraries

```
In [ ]: 1 ### Pandas and Numpy
2 import pandas as pd
3 import numpy as np
4
5 ### MongoDB Library
6 import pymongo
7
8 ### Machine Learning Libraries
9 from sklearn.model_selection import train_test_split
10 from sklearn.preprocessing import OneHotEncoder
11 from sklearn.preprocessing import StandardScaler
12 from sklearn.compose import make_column_transformer
13 from sklearn.linear_model import LogisticRegression
14 from sklearn.svm import SVC
15 from sklearn.model_selection import GridSearchCV
16 from sklearn.metrics import confusion_matrix, accuracy_score, classification
17
18 ### To ignore warnings
19 import warnings
20 warnings.filterwarnings('ignore')
```

2.0 Retrieving data from MongoDB

```
In [2]: 1 ### Retriving data from Mongoddb
2 ### creating connection with MongoDB
3
4 client = pymongo.MongoClient("mongodb+srv://{username}:{password}@clustershu
  < | >
```

```
In [3]: 1 db=client['Census_income']
2 collection=db['Census_income_data']
```

```
In [4]: 1 ### Locating our collection and data in MongoDB using find() method
2 data_from_mongodb=collection.find()
```

```
In [5]: 1 ### converting data from MongoDB to Dataframe in pandas
2 data_mongodb=pd.DataFrame(data_from_mongodb)
```

In [6]:

```
1 ### first 5 records in dataset
2 data_mongodb.head()
```

Out[6]:

	_id	index	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship
0	63635b0506652a035edadc4c	31	20	Private	266015	Some_college	10	Never_married	Sales	Unmarried
1	63635b0506652a035edadc2d	0	39	other	77516	Bachelors	13	Never_married	Adm_clerical	Not_married
2	63635b0506652a035edadc3a	13	32	Private	205019	other	12	Never_married	Sales	Not_married
3	63635b0506652a035edadc3d	16	25	other	176756	HS_grad	9	Never_married	other	Unmarried
4	63635b0506652a035edadc2e	1	50	other	83311	Bachelors	13	Married_civ_spouse	Exec_managerial	Married

In [7]:

```
1 ### dropping _id and index feature from dataset imported from MongoDB
2 data_mongodb.drop(['_id','index'], axis=1, inplace=True)
3 data_mongodb.head()
```

Out[7]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship
0	20	Private	266015	Some_college	10	Never_married	Sales	Unmarried
1	39	other	77516	Bachelors	13	Never_married	Adm_clerical	Not_married
2	32	Private	205019	other	12	Never_married	Sales	Not_married
3	25	other	176756	HS_grad	9	Never_married	other	Unmarried
4	50	other	83311	Bachelors	13	Married_civ_spouse	Exec_managerial	Married

3.0 Model and Evaluation

3.1 Separating Independent and Dependent features

In [8]:

```
1 ### Splitting data into independent feature dataframe and dependent feature
2 X=data_mongodb.iloc[:, :-1]
3 y=data_mongodb.iloc[:, -1]
4 X.head()
```

Out[8]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship
0	20	Private	266015	Some_college	10	Never_married	Sales	Unmarried
1	39	other	77516	Bachelors	13	Never_married	Adm_clerical	Not_married
2	32	Private	205019	other	12	Never_married	Sales	Not_married
3	25	other	176756	HS_grad	9	Never_married	other	Unmarried
4	50	other	83311	Bachelors	13	Married_civ_spouse	Exec_managerial	Married

In [9]: 1 y.head()

Out[9]: 0 0
1 0
2 0
3 0
4 0
Name: salary, dtype: int64

3.2 Train Test Split

In [10]: 1 *### random state train test split will be same with all people using random_*
2
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, ra
4 X_train.head()

Out[10]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation
34576	30	Private	23778	Some_college	10	Never_married	Exec_managerial
33148	41	Private	112763	Some_college	10	Married_civ_spouse	Adm_clerical
2109	20	Private	241752	HS_grad	9	Married_civ_spouse	other
33501	35	Private	211494	Bachelors	13	Never_married	Exec_managerial
47110	24	Private	408585	other	4	Married_civ_spouse	other

In [11]: 1 y_train.head()

Out[11]: 34576 0
33148 0
2109 0
33501 0
47110 0
Name: salary, dtype: int64

In [12]: 1 X_test.head()

Out[12]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	re
3436	52	Private	48925	Some_college	10	Married_civ_spouse	Adm_clerical	
16332	27	Private	31659	Bachelors	13	Married_civ_spouse	other	
39798	67	Private	187553	other	4	Divorced	Prof_specialty	Nc
12405	45	other	255559	HS_grad	9	Never_married	Adm_clerical	Nc
7584	32	Private	169955	Some_college	10	Married_civ_spouse	Other_service	

```
In [13]: 1 y_test.head()
```

```
Out[13]: 3436      0
          16332     1
          39798     0
          12405     0
          7584      0
          Name: salary, dtype: int64
```

```
In [14]: 1
          2 ### both will have same shape
          3 X_train.shape, y_train.shape
```

```
Out[14]: ((36609, 14), (36609,))
```

```
In [15]: 1 ### both will have same shape
          2 X_test.shape, y_test.shape
```

```
Out[15]: ((12204, 14), (12204,))
```

3.3 Feature Encoding

```
In [16]: 1 column_trans=make_column_transformer(
          2     (OneHotEncoder(), ['workclass', 'education', 'marital_status', 'occup
          3     remainder='passthrough'])
```

```
In [17]: 1 X_train=column_trans.fit_transform(X_train)
```

```
In [18]: 1 X_test=column_trans.transform(X_test)
```

3.4 Feature Scaling

```
In [19]: 1 scaler=StandardScaler()
```

```
In [20]: 1 X_train=scaler.fit_transform(X_train)
```

```
In [21]: 1 X_test=scaler.transform(X_test)
```

```
In [22]: 1 X_train_scaled=pd.DataFrame(X_train)
          2 X_train_scaled.head()
```

```
Out[22]:
```

	0	1	2	3	4	5	6	7	8
0	0.573765	-0.573765	-0.444065	-0.691388	1.869548	-0.638185	-0.394399	-0.923205	1.426647
1	0.573765	-0.573765	-0.444065	-0.691388	1.869548	-0.638185	-0.394399	1.083183	-0.700944
2	0.573765	-0.573765	-0.444065	1.446367	-0.534889	-0.638185	-0.394399	1.083183	-0.700944
3	0.573765	-0.573765	2.251920	-0.691388	-0.534889	-0.638185	-0.394399	-0.923205	1.426647
4	0.573765	-0.573765	-0.444065	-0.691388	-0.534889	1.566943	-0.394399	1.083183	-0.700944

5 rows × 34 columns

▢

◀ ▶

```
In [23]: 1 X_test_scaled=pd.DataFrame(X_test)
          2 X_test_scaled.head()
```

```
Out[23]:
```

	0	1	2	3	4	5	6	7	8
0	0.573765	-0.573765	-0.444065	-0.691388	1.869548	-0.638185	-0.394399	1.083183	-0.700944
1	0.573765	-0.573765	2.251920	-0.691388	-0.534889	-0.638185	-0.394399	1.083183	-0.700944
2	0.573765	-0.573765	-0.444065	-0.691388	-0.534889	1.566943	2.535503	-0.923205	-0.700944
3	-1.742874	1.742874	-0.444065	1.446367	-0.534889	-0.638185	-0.394399	-0.923205	1.426647
4	0.573765	-0.573765	-0.444065	-0.691388	1.869548	-0.638185	-0.394399	1.083183	-0.700944

5 rows × 34 columns

▢

◀ ▶

3.5 Logistic Regression Model

```
In [24]: 1 ### model
          2 logistic_reg=LogisticRegression()
          3 logistic_reg
```

```
Out[24]: LogisticRegression()
```

```
In [25]: 1 logistic_reg.fit(X_train, y_train)
```

```
Out[25]: LogisticRegression()
```

```
In [26]: 1 logistic_reg_pred=logistic_reg.predict(X_test)
          2 logistic_reg_pred
```

```
Out[26]: array([0, 1, 0, ..., 0, 0, 0])
```

```
In [27]: 1 confusion_mat=confusion_matrix(y_test, logistic_reg_pred)
         2 confusion_mat
```

```
Out[27]: array([[8618,  644],
               [1175, 1767]])
```

```
In [28]: 1 truly_positive=confusion_mat[0][0]
         2 falsely_positive=confusion_mat[0][1]
         3 falsely_negative=confusion_mat[1][0]
         4 truly_negative=confusion_mat[1][1]
```

```
In [29]: 1 classification_rep_log_reg=classification_report(y_test, logistic_reg_pred)
         2 print(classification_rep_log_reg)
```

	precision	recall	f1-score	support
0	0.88	0.93	0.90	9262
1	0.73	0.60	0.66	2942
accuracy			0.85	12204
macro avg	0.81	0.77	0.78	12204
weighted avg	0.84	0.85	0.85	12204

3.6 Support Vector Classifier Model

```
In [30]: 1 X_train1, X_test1, y_train1, y_test1 = train_test_split(X, y, test_size=0.25)
```

```
In [31]: 1 column_trans_svc=make_column_transformer(
         2     (OneHotEncoder(), ['workclass', 'education', 'marital_status', 'occup
         3     remainder='passthrough'])
```

```
In [32]: 1 X_train1=column_trans_svc.fit_transform(X_train1)
```

```
In [33]: 1 X_test1=column_trans_svc.transform(X_test1)
```

```
In [34]: 1 scaler_svc=StandardScaler()
         2 scaler_svc
```

```
Out[34]: StandardScaler()
```

```
In [35]: 1 X_train1=scaler_svc.fit_transform(X_train1)
```

```
In [36]: 1 X_test1=scaler_svc.transform(X_test1)
```

```
In [ ]: 1 svc=SVC()
        2 svc
```

Out[38]: SVC()

```
In [ ]: 1 svc.fit(X_train1, y_train1)
```

Out[39]: SVC()

```
In [ ]: 1 svc_pred=svc.predict(X_test1)
        2 svc_pred
```

Out[40]: array([0, 1, 0, ..., 0, 0, 0])

```
In [ ]: 1 confusion_mat_svc=confusion_matrix(y_test1, svc_pred)
        2 confusion_mat_svc
```

Out[41]: array([[8718, 544],
[1256, 1686]])

```
In [ ]: 1 truly_positive=confusion_mat_svc[0][0]
        2 falsely_positive=confusion_mat_svc[0][1]
        3 falsely_negative=confusion_mat_svc[1][0]
        4 truly_negative=confusion_mat_svc[1][1]
```

```
In [ ]: 1 classification_rep_svc=classification_report(y_test1, svc_pred)
        2 print(classification_rep_svc)
```

	precision	recall	f1-score	support
0	0.87	0.94	0.91	9262
1	0.76	0.57	0.65	2942
accuracy			0.85	12204
macro avg	0.82	0.76	0.78	12204
weighted avg	0.85	0.85	0.85	12204

3.7 Hyper-Parameter Tuning Logistic Regression Model

```
In [ ]: 1 param_grid = [
        2     {'penalty' : ['l1', 'l2', 'elasticnet', 'none'],
        3     'C' : np.logspace(-4, 4, 5),
        4     'solver' : ['lbfgs', 'newton-cg', 'liblinear', 'sag', 'saga'],
        5     'max_iter' : [100, 500]
        6     }
        7 ]
```

```
In [ ]: 1 log_reg_hpt=LogisticRegression()
        2 log_reg_hpt
```

Out[45]: LogisticRegression()


```
In [ ]: 1 hpt_log_reg=GridSearchCV(log_reg_hpt, param_grid = param_grid)
```

```
In [ ]: 1 best_hpt_log_reg=hpt_log_reg.fit(X_train, y_train)
2 best_hpt_log_reg
```

```
Out[47]: GridSearchCV(estimator=LogisticRegression(),
                      param_grid=[{'C': array([1.e-04, 1.e-02, 1.e+00, 1.e+02, 1.e+04]),
                                   'max_iter': [100, 500],
                                   'penalty': ['l1', 'l2', 'elasticnet', 'none'],
                                   'solver': ['lbfgs', 'newton-cg', 'liblinear', 'sag',
                                             'saga']}]])
```

```
In [ ]: 1 ### getting best parameters for Logistic Regression model after gridsearchCV
2 print("Best parameters are {} for optimal accuracy.".format(best_hpt_log_reg
```

Best parameters are LogisticRegression(C=0.01, penalty='l1', solver='liblinear') for optimal accuracy.

```
In [ ]: 1 ### getting best accuracy for Logistic Regression model after gridsearchCV
2 print("Best accuracy is {}".format(best_hpt_log_reg.score(X_test, y_test)))
```

Best accuracy is 0.8504588659455916

3.8 Hyper-Parameter Tuning Support Vector Classifier Model

```
In [37]: 1 svc_hpt=SVC()
2 svc_hpt
```

```
Out[37]: SVC()
```

```
In [38]: 1 #### using gridsearchcv to increase model efficiency by combining above para
2 param_grid={'C':[1,2,3], 'kernel':['rbf']}
3 hpt_svc=GridSearchCV(svc_hpt, param_grid=param_grid)
```

```
In [39]: 1 best_hpt_svc=hpt_svc.fit(X_train1, y_train1)
2 best_hpt_svc
```

```
Out[39]: GridSearchCV(estimator=SVC(),
                      param_grid={'C': [10, 20], 'degree': [2, 3],
                                   'kernel': ['linear', 'rbf', 'poly', 'sigmoid']})
```

```
In [40]: 1 ### getting best parameters for Logistic Regression model after gridsearchCV
2 print("Best parameters are {} for optimal accuracy.".format(best_hpt_svc.bes
```

Best parameters are SVC(C=20, degree=2, kernel='poly') for optimal accuracy.

```
In [41]: 1 ### getting best accuracy for Logistic Regression model after gridsearchCV
2 print("Best accuracy is {}".format(best_hpt_svc.score(X_test, y_test)))
```

Best accuracy is 0.8491478203867584

```
In [42]: 1 svc_hpt1=SVC()  
        2 svc_hpt1
```

Out[42]: SVC()

```
In [43]: 1 #### using gridsearchcv to increase model efficiency by combining above para  
        2 param_grid1={'C':[1,2,3], 'kernel':['rbf']}  
        3 hpt_svc1=GridSearchCV(svc_hpt1, param_grid=param_grid1)
```

```
In [44]: 1 best_hpt_svc1=hpt_svc1.fit(X_train1, y_train1)  
        2 best_hpt_svc1
```

Out[44]: GridSearchCV(estimator=SVC(), param_grid={'C': [1, 2, 3], 'kernel': ['rbf']})

```
In [45]: 1 ### getting best parameters for Logistic Regression model after gridsearchCV  
        2 print("Best parameters are {} for optimal accuracy.".format(best_hpt_svc1.be  
Best parameters are SVC(C=2) for optimal accuracy.
```

```
In [47]: 1 ### getting best accuracy for Logistic Regression model after gridsearchCV  
        2 print("Best accuracy is {}".format(best_hpt_svc1.score(X_test1, y_test1)))  
Best accuracy is 0.8517699115044248
```

Note: Please refer my github repo for ROC AUC curve implementation