

Machine Learning Basics

In this session I am practicing the models in Machine Learning using Python

Topics Covered :-

1. Linear Regression
2. Rigid and Lasso
3. Logistic Regression

#sklearn Linear Model Regression

Linear Regression,Rigid and Lasso

```
In [1]: 1 #house Pricing dataset Boston House
        2 from sklearn.datasets import load_boston
```

```
In [2]: 1 import numpy as np
        2 import pandas as pd
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 %matplotlib inline
        6
        7 ##importing all the required files
```

```
In [3]: 1 ## assigning to df
        2 df=load_boston()
```

```
In [4]: 1 ## to show the dataset
        2 df
```

```
Out[4]: {'data': array([[6.3200e-03, 1.8000e+01, 2.3100e+00, ..., 1.5300e+01, 3.9690e+02,
        4.9800e+00],
        [2.7310e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9690e+02,
        9.1400e+00],
        [2.7290e-02, 0.0000e+00, 7.0700e+00, ..., 1.7800e+01, 3.9283e+02,
        4.0300e+00],
        ...,
        [6.0760e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
        5.6400e+00],
        [1.0959e-01, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9345e+02,
        6.4800e+00],
        [4.7410e-02, 0.0000e+00, 1.1930e+01, ..., 2.1000e+01, 3.9690e+02,
        7.8800e+00]]),
        'target': array([24. , 21.6, 34.7, 33.4, 36.2, 28.7, 22.9, 27.1, 16.5, 18.9,
        15. ,
        18.9, 21.7, 20.4, 18.2, 19.9, 23.1, 17.5, 20.2, 18.2, 13.6, 19.6,
        15.2, 14.5, 15.6, 13.9, 16.6, 14.8, 18.4, 21. , 12.7, 14.5, 13.2,
        13.1, 13.5, 18.9, 20. , 21. , 24.7, 30.8, 34.9, 26.6, 25.3, 24.7,
        21.2, 10.2, 22. , 16.6, 11.4, 10.4, 10.7, 22.5, 25. , 22.4, 10.2])}
```

```
In [5]: 1 ## framed dataset
        2 pd.DataFrame(df.data)
```

```
Out[5]:
```

	0	1	2	3	4	5	6	7	8	9	10	11	12
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33
...
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0	273.0	21.0	396.90	7.88

506 rows × 13 columns

```
In [6]: 1 dataset=pd.DataFrame(df.data)
        2 dataset.columns=df.feature_names
        3 print(dataset.head())
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	

	PTRATIO	B	LSTAT
0	15.3	396.90	4.98
1	17.8	396.90	9.14
2	17.8	392.83	4.03
3	18.7	394.63	2.94
4	18.7	396.90	5.33

```
In [7]: 1 ## top 5 values
        2 dataset.head()
```

Out[7]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [8]: 1 ## targeting the price
        2 dataset['price']=df.target
```

```
In [9]: 1 dataset.head()
```

Out[9]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [10]: 1  ## Dividing the dataset into independent and dependent features
        2
        3  ##Indexing remving the Last column
        4  X=dataset.iloc[:, :-1] ##independent feature
        5
        6  y=dataset.iloc[:, -1]##dependent feature
```

```
In [11]: 1  X.head()
```

```
Out[11]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [12]: 1  y.head()
```

```
Out[12]: 0    24.0
         1    21.6
         2    34.7
         3    33.4
         4    36.2
Name: price, dtype: float64
```

```
In [13]: 1  ## Linear Regression
        2  from sklearn.linear_model import LinearRegression
        3  from sklearn.model_selection import cross_val_score
        4  lin_reg=LinearRegression()
        5  mse=cross_val_score(lin_reg,X,y,scoring='neg_mean_squared_error',cv=5) ##cro
        6  print(mse)
```

```
[-12.46030057 -26.04862111 -33.07413798 -80.76237112 -33.31360656]
```

```
In [14]: 1  mean_mse=np.mean(mse)
        2  print(mean_mse)
```

```
-37.13180746769922
```

```
lin_reg.predict(4) #to predict the value
```

```
In [15]: 1  ## Ridge Regression
```

```
In [21]: 1 from sklearn.linear_model import Ridge
2 from sklearn.model_selection import GridSearchCV
3 ridge=Ridge()
4 params={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-3,1e-2,1,5,10,20]}
5
6 ridge_regressor=GridSearchCV(ridge,params,scoring='neg_mean_squared_error',c
7 ridge_regressor.fit(X,y)
8
9
```

```
Out[21]: GridSearchCV(cv=5, estimator=Ridge(),
                    param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.001, 0.01, 1,
                    5, 10, 20]},
                    scoring='neg_mean_squared_error')
```

```
In [23]: 1 print(ridge_regressor.best_params_)
2 print(ridge_regressor.best_score_)
```

```
{'alpha': 20}
-32.38025025182513
```

```
In [24]: 1 from sklearn.linear_model import Lasso
2 from sklearn.model_selection import GridSearchCV
3 lasso=Lasso()
4 params={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-3,1e-2,1,5,10,20]}
5
6 lasso_regressor=GridSearchCV(lasso,params,scoring='neg_mean_squared_error',c
7 lasso_regressor.fit(X,y)
8
9
```

C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 4430.746729651311, tolerance: 3.9191485420792076

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 4397.459304778431, tolerance: 3.3071316790123455
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 3796.653037433508, tolerance: 2.813643886419753
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 2564.292735790545, tolerance: 3.3071762123456794
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 4294.252997826028, tolerance: 3.4809104444444445
```

```
model = cd_fast.enet_coordinate_descent(
```

```
Out[24]: GridSearchCV(cv=5, estimator=Lasso(),
                  param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.001, 0.01, 1,
                                         5, 10, 20]},
                  scoring='neg_mean_squared_error')
```

```
In [25]: 1 print(lasso_regressor.best_params_)
2 print(lasso_regressor.best_score_)
```

```
{'alpha': 1}
-35.531580220694856
```

```
In [26]: 1 from sklearn.linear_model import Ridge
2 from sklearn.model_selection import GridSearchCV
3 ridge=Ridge()
4 params={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-3,1e-2,1,5,10,20,30,40,45,50,55,100]}
5
6 ridge_regressor=GridSearchCV(ridge,params,scoring='neg_mean_squared_error',c
7 ridge_regressor.fit(X,y)
8
9
```

```
Out[26]: GridSearchCV(cv=5, estimator=Ridge(),
                    param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.001, 0.01, 1,
                    5, 10, 20, 30, 40, 45, 50, 55, 100]}},
                    scoring='neg_mean_squared_error')
```

```
In [27]: 1 from sklearn.linear_model import Lasso
2 from sklearn.model_selection import GridSearchCV
3 lasso=Lasso()
4 params={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-3,1e-2,1,5,10,20,30,40,45,50,55,100]}
5
6 lasso_regressor=GridSearchCV(lasso,params,scoring='neg_mean_squared_error',c
7 lasso_regressor.fit(X,y)
8
9
```

```
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 4430.746729651311, tolerance: 3.9191485420792076
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 4397.459304778431, tolerance: 3.3071316790123455
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 3796.653037433508, tolerance: 2.813643886419753
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 2564.292735790545, tolerance: 3.3071762123456794
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\arya shriva\anaconda3\lib\site-packages\sklearn\linear_model\_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 4294.252997826028, tolerance: 3.4809104444444445
```

```
model = cd_fast.enet_coordinate_descent(
```

```
Out[27]: GridSearchCV(cv=5, estimator=Lasso(),
                    param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.001, 0.01, 1,
                    5, 10, 20, 30, 40, 45, 50, 55, 100]}},
                    scoring='neg_mean_squared_error')
```

```
In [28]: 1 print(lasso_regressor.best_params_)
2 print(lasso_regressor.best_score_)
```

```
{'alpha': 1}
-35.531580220694856
```

```
In [30]: 1 from sklearn.model_selection import train_test_split
2 X_train, X_test, y_train, y_test = train_test_split(
3 X, y, test_size=0.33, random_state=42)
```

```
In [31]: 1 #ridge_regressor
2 from sklearn.linear_model import Ridge
3 from sklearn.model_selection import GridSearchCV
4 ridge=Ridge()
5 params={'alpha':[1e-15,1e-10,1e-8,1e-3,1e-3,1e-2,1,5,10,20,30,40,45,50,55,100]}
6
7 ridge_regressor=GridSearchCV(ridge,params,scoring='neg_mean_squared_error',cv=5)
8 ridge_regressor.fit(X,y)
```

```
Out[31]: GridSearchCV(cv=5, estimator=Ridge(),
                    param_grid={'alpha': [1e-15, 1e-10, 1e-08, 0.001, 0.001, 0.01, 1,
                                           5, 10, 20, 30, 40, 45, 50, 55, 100]},
                    scoring='neg_mean_squared_error')
```